

Quantifying Differential Methylation in Structured Populations

Önder Kartal

August 31, 2016

Abstract

The mapping and quantification of epigenetic differences between cells, tissues, and organisms at genome scale is critical to determine effects of environment, differentiation, and reproduction on gene function. The techniques to detect epigenetic signals are based on next-generation sequencing. However, this technology only produces a random sample of sequence reads and a crucial part of signal detection is done in silico. In the case of DNA methylation, sodium bisulfite treated reads are aligned to the reference genome to build genome-scale methylation tables.

Here, we address the problem of reliably detecting differentially methylated cytosines (DMCs) based on methylation tables. We investigate an information-theoretic measure, the Jensen-Shannon Divergence (JSD), to scan the genome for DMCs. JSD is a well-founded and very general measure of statistical dependence that can be applied to any probabilistic description of the methylation state and does not suffer from shortcomings of correlation-based measures or standard statistical tests.

We set up a computational pipeline to calculate JSD along the genome for a population. We extend the definition of JSD to account for population structure similar to Sewall Wright's F-statistics. This allows the detection of DMCs between subpopulations, for example different lines, demes, or different treatment groups. The analysis of published methylomes of *Arabidopsis* reveals that population diversity is consistently higher for cytosines in a CG context than those in a CHG or CHH context. As CG methylation is enriched in gene bodies, especially exons, this suggests differences in gene activity.

In conclusion, we have established that JSD allows the straightforward and accurate detection of DMCs while additionally quantifying the extent of differentiation.

1 Introduction

In recent years it has become increasingly clear that gene function is not determined by genetic information, that is DNA sequence, alone. Epigenetics is

the study of mitotically and/or meiotically heritable changes in gene function that can not be explained by changes in DNA sequence CITE

Cytosine (C) Methylation is one of the best-characterized epigenetic signals. It is associated with gene silencing when found in promoters and with active

transcription when found in gene bodies. It is also an effective mechanism to silence transposons.

The state-of-the-art method to measure the methylation state across the whole genome at single-base-pair resolution is to sequence bisulfite-treated DNA fragments. This is followed essentially by two computational steps CITE: read alignment to the reference genome and inference of the methylation level at each locus. This analysis results in a methylation table that allows to retrieve the number of methylated and unmethylated reads for each covered C.

The possibility to generate methylomes with relative ease has led to large scale studies that try to investigate the extent of differential methylation in a population of cells or organisms. By differential methylation we mean single differentially methylated cytosines (DMCs) or spatially clustered collections of DMCs, called differentially methylated regions (DMRs). Typically, the sample population is structured in the sense that each individual belongs at least to a certain group. The partitioning can be, for example, in terms of lineages, generations, demes or experimental treatments.

A basic task is then to quantify the methylation diversity across the genome within and between subpopulations. We argue that the most reasonable way to do so, is to use an information-theoretic measure based on entropy. Our basic reasoning goes as follows: Our degree of belief regarding the methylation state at a locus is described by a probability distribution function (PDF). This representation of our knowledge accounts as well for the uncertainty in a quantitative fashion. In this setting, assessing the difference between loci means measuring the difference between PDFs. Information theory provides the appropriate measures to quantify the difference between PDFs by quantifying how much information is shared between the PDFs. This is the most general measure of statistical dependence we are aware of. In particular, the Jensen-Shannon divergence is a suitable measure of differences in a population of PDFs.

2 Material and Methods

Our method of inference is based on comparing probabilistic representations of methylation states. To follow the argument, it is necessary to understand central concepts from probability theory and information theory. Here, we want to give a brief overview of these concepts and their implementation in computer code to scan genomic regions based on methylome data.

2.1 Basic Concepts

A probabilistic model requires the definition of a **sample space** Ω . This set contains all possible states of a system or at least those states that we care to model. If we conduct an experiment we can predict the outcome only to a certain degree due to missing information. The uncertainty is either about the laws that govern the system, about parameter values that are outside of our control or about all the intricacies of the experiment itself. Another common source of uncertainty is the deliberate ignorance of processes that may interfere with the modelled process but whose influence we are unable to specify. Therefore, we model such an experiment as a **random variable** X . The name is misleading, since X is actually a function, $X : \Omega \mapsto \mathbb{R}$, that maps each possible state to

a real number. Thus, it is customary to also refer to the value of X as the random variable. We are not keen, at the moment, to resolve this ambiguity in the literature. In essence, the value of a random variable is a proxy for a single or several states of the system and its observation is an **event**. A random variable that maps a non-singleton subset of its domain Ω to a single value is typically used if we are interested in an outcome that is compatible with any of the states in the subset. An event written as $(X = x)$, where $x \in X(\Omega)$, is defined formally in terms of the sample space as

$$(X = x) = \{\omega : X(\omega) = x\}. \quad (1)$$

A **probability measure** \Pr is a function that assigns a number in the range $[0, 1]$ to each event and fulfills the condition $\Pr(\Omega) = 1$. The normalization guarantees that probability is a conserved quantity that can only be redistributed among the events. The notion of a probability measure is closely related to that of a **probability distribution**. The distribution function is central to our approach to detect differences in methylation. In the discrete case, the probability is concentrated in a countable set of points and the discrete distribution is called a **probability mass function** (pmf). For N possible outcomes, $X = \{x_1, x_2, \dots, x_N\}$, the probabilities of the events are given by $\Pr(X = x_1)$, $\Pr(X = x_2)$, etc., and the normalization condition reads $\sum_k \Pr(X = x_k) = 1$. We will denote the pmf as $p(x_k)$ instead of the more verbose $\Pr(X = x_k)$. If the random variable is continuous, an event is represented by an interval $(a < X < b)$ whose probability is

$$\Pr(a < X < b) = \int_a^b p(x) dx. \quad (2)$$

The normalization condition reads $\Pr(\Omega) = \int_{X(\Omega)} p(x) dx = 1$, where $p(x)$ is the **probability density function** (pdf). In contrast to the pmf, note that $p(x) \neq \Pr(X = x)$.

If we consider more than one random variable or a single event we are concerned with questions like: How likely is it that distinct events will co-occur? Does the probability of an event change if we have information that another event applies? What does it mean that two events are independent? To answer these questions we have to introduce the important concept of **conditional probability**. When we write $\Pr(A)$ for the probability of any event we actually mean $\Pr(A | I)$, where the notation ‘|’ means ‘conditional on’ or ‘given that’. The symbol I summarizes all background information that is implied when we speak about probabilities of A . Especially in the Bayesian camp, conditional probabilities are often taken as the primitive notions of probability theory (Jaynes, 2003; Sivia and Skilling, 2006). According to this viewpoint, there is no such thing as an unconditional probability. Based on this qualification, the product rule explains how **joint probabilities** are composed. For two events A and B we have

$$\begin{aligned} \Pr(AB | I) &= \Pr(A | BI) \times \Pr(B | I) \\ &= \Pr(B | AI) \times \Pr(A | I), \end{aligned} \quad (3)$$

The conditional probabilities on the right-hand side highlight the fact that the events may be dependent. If the two events are independent then conditioning

on one event does not confer any information about the other event and we can write $\Pr(A | BI) = \Pr(A | I)$ (analogous for $\Pr(B | AI)$). This leads to the familiar equation for **statistical independence**,

$$\Pr(AB | I) = \Pr(A | I) \times \Pr(B | I). \quad (4)$$

The existence of joint probabilities carries over to distribution functions, of course. For two random variables X and Y we will write the joint distributions as $p(x_k, y_k)$ for the pmf and as $p(x, y)$ for the pdf.

Two central notions of this paper are Shannon Entropy and Jensen-Shannon Divergence. Claude Shannon noticed that a quantity he called entropy (now commonly referred to as information or Shannon entropy) can be used to define information rigorously and to develop a mathematical theory of communication (Shannon, 1948).

Definition 1 (Shannon Entropy). *Given a discrete PDF f the Shannon entropy H quantifies the average degree of uncertainty in predicting any elementary event $x_k \in X$:*

$$H(f) = \sum_k p(x_k) \log_2 \frac{1}{p(x_k)} \quad (\text{in bits}). \quad (5)$$

The Jensen-Shannon divergence was explicitly introduced as a divergence measure to quantify the similarity of PDFs by Lin (1991).

Definition 2 (Jensen-Shannon Divergence). *For a set of distributions $\{f\} = \{f_1, f_2, \dots\}$ defined over the same sample space X and weighted by factors $\pi_i \geq 0$ with $\sum_i \pi_i = 1$ the Jensen-Shannon divergence is*

$$D_{JS}(\{f\}) = H\left(\sum_i \pi_i f_i\right) - \sum_i \pi_i H(f_i), \quad (\text{in bits}). \quad (6)$$

Note, that on the left hand side of Eq. 6 we have the entropy of the finite mixture of the input distributions, $H(f_{\text{mix}})$, from which the (weighted) average entropy \bar{H} of the distributions is subtracted. Figure ?? illustrates that this difference is always non-negative and thus mixing of distributions can only lead to an increase in entropy, hence uncertainty. In short, JSD quantifies the difference between distributions by quantifying the information loss due to mixing.

This is also evident by rewriting D_{JS} in terms of Kullback-Leibler Divergence (KLD), D_{KL} ,

$$D_{JS} = \sum_i \pi_i D_{KL}(f_i || f_{\text{mix}}). \quad (7)$$

The KLD $D_{KL}(S || T)$ is a directed (non-symmetric) measure of information loss when approximating a given source distribution S by a target distribution T Kullback and Leibler, 1951. By using the mixture as the target distribution and averaging over the distances, JSD becomes a smoothed and symmetric version of KLD. We will discuss how JSD is also a measure of statistical dependence and its close relation with both mutual information and Fisher's information metric.

Diversity vs. diversity indices...

fig/jsd_entropy-eps-converted-to.pdf

Figure 1: Jensen-Shannon Divergence as mixing entropy. The graph shows the entropy H of a discrete binary probability distribution function (PDF) $f(\{x_1, x_2\})$ as a function of the probability $p = f(x_1) = 1 - f(x_2)$. The black points show the location of three example PDFs, f_1 , f_2 and f_3 , in this space. Features highlighted in red are quantities entering the Jensen-Shannon Divergence (JSD). The entropy of the mixture $H(f_{\text{mix}})$ can lie anywhere on the red segment of the curve. The average entropy, $\langle H \rangle$, is restricted to the area designated by the triangle.

2.2 Implementation

2.3 Data Sources

3 Results and Discussion

3.1 Jensen-Shannon Divergence in Structured Populations

It is often necessary or desirable to quantify differences between populations since most natural and experimental populations are structured (or subdivided) inherently due to genetic sampling or can be subdivided into groups that arise

by design of the experiment. The subdivision into groups can be performed at different levels if hierarchies are present in the population, see Fig. ?? . At any level, it becomes necessary to consider within-group divergences in order to derive between-group divergences. As it stands, the expression for JSD gives the within-group divergence only. This partitioning of divergence is in principle reminiscent of the partitioning of variance in ANOVA to calculate differences between two groups.

Illustration of Population structure.

The notion of population structure can be made precise in terms of equivalence relations. An equivalence relation R over a set T (the population) induces a partition of the set into subsets called equivalence classes (usually referred to as subgroups). Any two members of the same equivalence class, a and b , are said to be R -equivalent or equivalent by R , written as $a \sim_R b$. The set of all equivalence classes is called the quotient set of T by R , denoted by T/R . Of course, different equivalence relations can be defined on one and the same set, which would lead to different partitions. In particular, two different equivalence relations, say Q and R , can induce partitions at different levels in the sense that $a \sim_Q b$ implies $a \sim_R b$. Q is then said to be a finer relation than R and the Q -equivalence class of a member a is a subset of its R -equivalence class.

We give an example that is relevant for populations of Arabidopsis accessions. The relation K =continent would lead to the quotient set T/K that contains equivalence classes combining samples originating from the same continent. The relation L =country is a finer relation than K that leads to equivalence classes by countries in T/L . Obviously, for any given member a , its L -equivalence class is a subset of its K -equivalence class, since accessions from the same country are also from the same continent of origin.

Using the concept and terminology of equivalence relations, the between-group divergence in a set T partitioned by R is actually the divergence of the quotient set, $J(T/R)$. We want to show how to compute this quantity and how it is related to the total divergence $J(T)$ and the within-group divergences $J(S)$, $S \in T/R$,

$$J(T) = H(\pi_i^T p^i) - \pi_i^T H(p^i) \quad \text{and} \quad (8)$$

$$J(S) = H(\pi_i^S p^i) - \pi_i^S H(p^i). \quad (9)$$

The individual weight vectors π^S and π^T have the same size but for any subset S we have $\sum_j \pi_j^S [j \in S] = 1$ and $\pi_i^S = 0$ for $i \notin S$. Due to the convexity of entropy, the average divergence over the quotient set can never exceed the total divergence,

$$J(T) \geq \overline{J(S)} \quad \text{with} \quad \overline{J(S)} = \sum_s \pi_s^T J(s) [s \in T/R], \quad (10)$$

where the weights π_s^T now refer to subsets rather than individual members of T . It turns out that the mismatch between total and average divergence is precisely due to the divergence *between* the subgroups, $J(T/R)$:

$$J(T/R) = J(T) - \overline{J(S)} \quad (11)$$

$$= H(\pi_i^T p^i) - \pi_s^T H(\pi_i^s p^i), \quad (12)$$

which follows from the identity $\pi_i^T = \pi_S^T \pi_i^S$ for any given set S and the canceling of the term $\pi_i^T H(p^i)$.

Discuss the results, relation to Fst, alpha/beta/gamma Entropy, anova.

Sewall Wright used the generic notation ST for between-group correlation in his F -statistics Wright, 1949. We believe that using the notation of equivalence relations is more precise and especially more explicit since it is immediately clear by which relation the reference set is partitioned.

3.2 Differentially Methylated Positions

In this section we demonstrate that JSD allows to scan the genome of a species for differentially methylated positions (DMPs) in an arbitrarily subdivided population. For this we assume that a statistical sample of n individuals has been taken from a population. We further assume that the methylation state at each cytosine position for each individual is described in terms of a PDF. Population subdivision refers to a classification of the individuals into (possibly nested) groups that share some common property or history. For example, a subgroup may consist of plants that come from the same habitat or just shared the same experimental treatment. In cell populations, individuals coming from the same tissue may form a group that is nested within a group at the level of organs.

To compare cytosine bases across a population sample using PDFs, a hypothesis space needs to be defined. The choice of a hypothesis space is not a trivial task – it already implies a model of the system and the data-generating process.

3.2.1 DMPs from discrete representations of methylation state

Model 1 (Bernoulli Distribution). *We assume two mutually exclusive propositions or hypotheses for an arbitrary but fixed locus:*

1. $m \equiv$ *The locus is methylated.*
2. $\neg m \equiv$ *The locus is not methylated.*

In this case the PDF $p(M|I)$ over the set $M = \{m, \neg m\}$ is given by the probabilities

$$p = \Pr(m|I), \tag{13}$$

$$q = \Pr(\neg m|I) = 1 - p. \tag{14}$$

The Bernoulli distribution is one possible model if we assume that the locus is either completely methylated or completely unmethylated. This condition is fulfilled for haploid cells with completely operative maintenance methylation or, to a certain extent, for cytosine bases in CG context for plants. It is not an appropriate model if allele-specific methylation or hemimethylation is possible.

If the background information I does not favor any of the two propositions, for example, we would encode our indifference by assigning the prior probabilities

$$p = q = \frac{1}{2}. \tag{15}$$

Similar to the case of a fair (i.e. unbiased) coin, this is a reasonable assignment if we are completely uncertain as to the methylation state of the locus. To perform the genomic scan, however, we are interested in the post-data PDFs, $p(M|I, D)$, where the data consists of bases that map to the locus in question. For that we interpret the whole sample of bases that map to the locus, with k methylated Cs out of N total reads, as a single Bernoulli trial. In that case, the posterior probability of methylation is equal to the sample mean

$$\Pr(m|I, D) = \frac{k}{N}. \quad (16)$$

That is, under the assumption of a Bernoulli likelihood function, we can use the frequency distribution as the probability distribution. This is a very crude estimate of the methylation probability without taking into account sampling bias that is especially problematic for low sample sizes. Random reads from a DNA library show uneven distribution or even gaps when mapped to the reference genome. However, frequency distributions are reasonable first approximations and have been successfully used previously in comparing sequences with JSD by Grosse et al., 2002. In fact, they have shown that a natural choice of the weights that enter Eq. 6 allow to mitigate the effect of sampling bias when comparing PDFs. Using simulations, they have shown, that sampling bias can be taken into account by using

$$\pi_i = \frac{N_i}{\sum_i N_i}. \quad (17)$$

3.2.2 DMPs from continuous representations of methylation state

Rather than taking the whole sample of reads as a Bernoulli trial as in 1 each single mapped base can be seen as a Bernoulli trial. This leads naturally to a binomial distribution.

Model 2 (Binomial Distribution). *We assume that the methylation bias lies between 0 and 1...*

The convenient choice in this case for the prior is that of beta distribution. In that case one speaks of a conjugate prior since the posterior PDF is also a beta distribution. It can be shown that the sample mean is the mode of the posterior PDF if we choose a binomial likelihood function with a flat prior. However, the posterior PDF now conveys additional information about the uncertainty of our estimate. In formulating our problem in a more sensible way we did not throw away the information that is contained in the sample size...

3.2.3 Discussion

3.3 Differentially Methylated Regions

3.3.1 JSD in annotated regions

3.3.2 JSD-based methylome segmentation

3.3.3 Discussion

4 Conclusions

References

- Grosse, Ivo et al. (2002). “Analysis of symbolic sequences using the Jensen-Shannon divergence.” eng. In: *Phys Rev E Stat Nonlin Soft Matter Phys* 65.4 Pt 1, p. 041905.
- Jaynes, E. T. (2003). *Probability Theory: The Logic of Science*. Cambridge University Press.
- Kullback, S. and R. A. Leibler (1951). “On Information and Sufficiency”. In: *The Annals of Mathematical Statistics* 22.1, pp. 79–86. ISSN: 00034851.
- Lin, J. (1991). “Divergence measures based on the Shannon Entropy”. In: *IEEE T. Inform. Theory*. 37, pp. 145–151.
- Shannon, C. E. (1948). “A mathematical theory of communication”. In: *The Bell System Technical Journal* 27, pp. 379–423.
- Sivia, D. S. and J. Skilling (2006). *Data analysis: a bayesian tutorial*. Oxford: Oxford University Press.
- Wright, Sewall (1949). “The genetical structure of populations”. In: *Ann. Hum. Genet.* 15.1, pp. 323–354. ISSN: 1469-1809.